

Exercise 2.2 Bandit Example

Lenin G. Falconi*

2025-03-31

1 A 4 armed bandit problem

Exercise 2.2: Bandit example Consider a k -armed bandit problem with $k = 4$ actions, denoted 1, 2, 3, and 4. Consider applying to this problem a bandit algorithm using ϵ -greedy action selection, sample-average action-value estimates, and initial estimates of $Q_1(a) = 0$, for all a . Suppose the initial sequence of actions and rewards is $A_1 = 1, R_1 = -1, A_2 = 2, R_2 = 1, A_3 = 2, R_3 = -2, A_4 = 2, R_4 = 2, A_5 = 3, R_5 = 0$. On some of these time steps the ϵ case may have occurred, causing an action to be selected at random. On which time steps did this definitely occur? On which time steps could this possibly have occurred?

A k arm bandit algorithm is:

Algorithm 1 ϵ -Greedy k -Armed Bandit Algorithm

```
1: Input: Action set  $A = \{1, 2, \dots, k\}$ , exploration rate  $\epsilon$ 
2: Initialize for all  $a \in A$ :  $Q_1(a) \leftarrow 0$  and count  $N_1(a) \leftarrow 0$ 
3: for  $t = 1, 2, \dots$  do
4:   if a random number in  $[0, 1]$  is less than  $\epsilon$  then
5:     Choose a random action  $a_t \in A$ 
6:   else
7:     Select  $A_t = \underset{a}{\operatorname{argmax}}(Q_t(a))$ 
8:   end if
9:   Execute action  $a_t$ , observe reward  $R_t$ 
10:  Increment:  $N_{t+1}(a_t) \leftarrow N_t(a_t) + 1$ 
11:  Update:  $Q_{t+1}(a_t) \leftarrow Q_t(a_t) + \frac{1}{N_t(a_t)}(R_t - Q_t(a_t))$ 
12: end for
```

*lenin.falconi@epn.edu.ec

1.1 Analyzing Greedy

Equation 2.3 in book has no subscripts but n seems to mean time step or step and the equation seems to be about a single state. Because of that I make those adjustments in the equation in the algorithm. Is that correct? The problem does not define any ϵ .

$t = 1, a = 1$:

$$N_1 = [0, 0, 0, 0]$$

$$Q_1 = [0, 0, 0, 0]$$

$$A = [1, 2, 3, 4]$$

$$A_1 = 1, R_1 = -1$$

$$N_1(1) = 0 + 1$$

$$Q_2(1) = Q_1(1) + \frac{1}{N_1(1)}[R_1 - Q_1(1)]$$

$$Q_2(1) = 0 + [-1]$$

$$Q_2(1) = -1$$

for $t = 1, a \in \{2, 3, 4\}$: $N_1(2) = N_1(3) = N_1(4) = 0$ Because none of those actions is taken. As a consequence $Q_1(2) = Q_1(3) = Q_1(4) = 0$. And the Q_1 vector would be:

$$Q_1 = [-1, 0, 0, 0]$$

So we continue with $t = 2, a = 2$. I will use to index the values of the vector starting in 1 $t = 2, a = 2$:

$$N_2 = [1, 0, 0, 0]$$

$$Q_2 = [-1, 0, 0, 0]$$

Getting $\text{argmax}(Q_1) \neq 1$

$$A = [1, 2, 3, 4]^a$$

$$A_1 = 2, R_1 = 1$$

$$N_2(2) = 0 + 1$$

$$Q_3(2) = Q_2(2) + \frac{1}{N_2(2)}[R_2 - Q_2(2)]$$

$$Q_3(2) = 0 + [1]$$

$$Q_3(2) = 1$$

In the same way as before, I am reasoning that $Q_2(3) = Q_2(4) = 0$. As a consequence for $t = 3$, the agent chooses $a = 2$ $t = 3, a = 2$:

$$N_3 = [1, 1, 0, 0]$$

$$Q_3 = [-1, 1, 0, 0]$$

Getting $\text{argmax}(Q_1) \neq 1$

$$A = [1, 2, 3, 4]^a$$

$$A_1 = 2, R_1 = 1$$

$$N_3(2) = 1 + 1$$

$$Q_4(2) = Q_3(2) + \frac{1}{N_3(2)}[R_3 - Q_3(2)]$$

$$Q_4(2) = 1 + 1/2 * (-2 - 1)$$

$$Q_4(2) = -1/2$$

Writing a python program. As I understand the problem, the idea is to estimate when an ϵ - greedy decision may have occurred given the sequence of states. Should I try to see what happens if I force to write the *bandit program* assuming an ϵ value. But in that case, do I need the other sequences or can they be ignored?

```
import numpy as np
A = [1,2,3,4]
Q = np.array([0,0,0,0], dtype=float)
N = np.array([0,0,0,0])
T = np.arange(6)
epsilon = 0.1
sequence = [(1,-1),(2,1),(2,-2), (2,2), (3,0)]
for n, (action, reward) in enumerate(sequence):
    print(f"step: {n+1}")
    print(f"current action: {action}")
    N[action-1]+=1
    print(f"N:{N}")
    Q[action-1] += 1/N[action-1]*(reward-Q[action-1])
    print(f"Q_{n+1+1}:{Q}")
    print("=====")
print(f"N:{N}")
print(f"Q:{Q}")
```

```
step: 1
current action: 1
N:[1 0 0 0]
Q_2:[-1.  0.  0.  0.]
=====
step: 2
current action: 2
N:[1 1 0 0]
Q_3:[-1.  1.  0.  0.]
=====
step: 3
current action: 2
N:[1 2 0 0]
Q_4:[-1. -0.5  0.  0.]
=====
```

```

step: 4
current action: 2
N:[1 3 0 0]
Q_5:[-1.          0.33333333  0.          0.          ]
=====
step: 5
current action: 3
N:[1 3 1 0]
Q_6:[-1.          0.33333333  0.          0.          ]
=====
N:[1 3 1 0]
Q:[-1.          0.33333333  0.          0.          ]

```

The first step gets a Q-value negative. Then, the argmax result could be either actions 2, 3 and 4, that have a value of 0. To break the tie, the agent chooses action 2. The ϵ case happens again in step 2, because $Q_3(2) = 1$. I think step 3 remains greedy choosing 2, but the result of the prediction for next state gives a negative value. A greedy approach would recommend actions 3 or 4 for step 4. Here greedy did not occur since for step 4 again action 2 is taken. Step 5 does not use greedy since $a_5 = \underset{a}{\operatorname{argmax}}(Q_5) = 2$.

1.2 Sample-averaged action value estimates

In this case I think we have to apply the equation:

$$Q_t(a) = \frac{\sum_{t=1}^{t-1} R_t \cdot \mathbb{1}_{A_t=a}}{\sum_{t=1}^{t-1} \mathbb{1}_{A_t=a}}$$

In this case I think we just calculate the product of Reward and Action for each action of the sequence:

$$\begin{aligned}
Q_t(1) &= \frac{\sum_{t=1}^5 R_t \cdot \mathbb{1}_{A_t=1}}{\sum_{t=1}^5 \mathbb{1}_{A_t=1}} = \frac{1}{1} \\
Q_t(2) &= \frac{\sum_{t=1}^5 R_t \cdot \mathbb{1}_{A_t=2}}{\sum_{t=1}^5 \mathbb{1}_{A_t=2}} = \frac{0 + 2(1) + 2(-2) + 2(2)}{3} = \frac{2}{3} \\
Q_t(3) &= \frac{\sum_{t=1}^5 R_t \cdot \mathbb{1}_{A_t=3}}{\sum_{t=1}^5 \mathbb{1}_{A_t=3}} = \frac{0 + 0 + 0 + 0 + 3(0)}{1} = 0
\end{aligned}$$

As a consequence $Q = [-1, 0.666, 0, 0]$

1.3 ϵ - Greedy Program

```

import numpy as np
A = [1,2,3,4]
Q = np.array([0,0,0,0], dtype=float)
N = np.array([0,0,0,0])
T = np.arange(6)
epsilon = 0.1

```

```

def bandit(action):
    sequence = [(1,-1),(2,1),(2,-2), (2,2), (3,0), (4,0)]
    rewards = [reward for act, reward in sequence if act == action]
    if not rewards:
        raise ValueError(f"No reward for action {action}")
    return np.random.choice(rewards)

for n in T:
    roll_dice = np.random.random()
    print(f"roll_dice: {roll_dice}")
    if roll_dice < epsilon:
        # explore
        action = np.random.randint(1,5)
        print(f"action explore:{action}")

    else:
        # exploit
        action = np.argmax(Q)+1
        print(f"action exploit:{action}")

    reward = bandit(action)
    print(f"step: {n+1}")
    print(f"current action: {action}")
    N[action-1]+=1
    print(f"N:{N}")
    Q[action-1] += 1/N[action-1]*(reward-Q[action-1])
    print(f"Q_{n+1+1}:{Q}")
    print("=====")
print(f"N:{N}")
print(f"Q:{Q}")

```

```

roll_dice: 0.41547343809500925
action exploit:1
step: 1
current action: 1
N:[1 0 0 0]
Q_2:[-1.  0.  0.  0.]
=====
roll_dice: 0.9297839425270602
action exploit:2
step: 2
current action: 2
N:[1 1 0 0]
Q_3:[-1.  2.  0.  0.]
=====

```

```
roll_dice: 0.7914951823916979
action exploit:2
step: 3
current action: 2
N:[1 2 0 0]
Q_4:[-1.  2.  0.  0.]
=====
roll_dice: 0.40242581791537846
action exploit:2
step: 4
current action: 2
N:[1 3 0 0]
Q_5:[-1.  2.  0.  0.]
=====
roll_dice: 0.3849103661520379
action exploit:2
step: 5
current action: 2
N:[1 4 0 0]
Q_6:[-1.  2.  0.  0.]
=====
roll_dice: 0.07361910965678975
action explore:3
step: 6
current action: 3
N:[1 4 1 0]
Q_7:[-1.  2.  0.  0.]
=====
N:[1 4 1 0]
Q:[-1.  2.  0.  0.]
```